

# **The Ethical Frontier of Artificial Intelligence: Navigating Societal Challenges and Charting a Responsible Path Forward**

## **I. Executive Summary**

Artificial intelligence (AI) stands as a transformative force, rapidly reshaping global industries and daily life with promises of unprecedented efficiency and problem-solving capabilities. However, this profound power is intrinsically linked to a complex array of ethical dilemmas that demand urgent and thoughtful consideration. This report provides a comprehensive examination of five critical ethical concerns associated with AI: Bias and Discrimination, Autonomy and Accountability, Job Displacement, Privacy and Surveillance, and AI in Weapons Development. Each domain presents unique challenges, yet they are deeply interconnected, forming a complex web where addressing one issue often has ripple effects on others. The analysis underscores the imperative for robust ethical frameworks, collaborative multi-stakeholder engagement, and adaptive regulatory mechanisms. The overarching conclusion is that AI's dual nature, encompassing both immense potential and significant risks, necessitates proactive, interdisciplinary approaches to governance to ensure its development and deployment align with fundamental human values and societal well-being.

## **II. Introduction: The Dual Nature of AI and the Imperative for Ethical Governance**

Artificial intelligence is rapidly reshaping industries and daily life, promising unprecedented efficiency and problem-solving capabilities.<sup>1</sup> This transformative power, however, is accompanied by a complex array of ethical dilemmas that demand careful consideration and proactive governance.<sup>1</sup> The ethical implications of AI are not merely theoretical; they have profound real-world repercussions across various sectors, including healthcare, finance, communication, and criminal justice.<sup>2</sup>

This report delves into five critical areas of AI ethics: Bias and Discrimination, Autonomy and Accountability, Job Displacement, Privacy and Surveillance, and AI in

Weapons Development. Each of these areas presents unique challenges and requires specific policy and technical interventions.

Ethical considerations in AI are not isolated but are deeply interdependent. For instance, algorithmic biases can lead to discriminatory outcomes, which in turn raise questions of accountability and may necessitate increased surveillance for mitigation, thereby impacting privacy. Similarly, job displacement can exacerbate societal inequalities, which biased AI systems might then amplify. Addressing these issues requires a holistic approach that recognizes these intricate relationships.<sup>1</sup>

A closer examination reveals a cyclical relationship between different ethical concerns. Bias in training data, a primary source of AI bias, can lead to discriminatory outcomes. These outcomes then necessitate robust accountability frameworks to assign responsibility for the harm caused. In response to such harms, there might be a perceived need for increased monitoring or surveillance to detect misuse or errors in AI systems. This, however, immediately raises new privacy concerns. If the surveillance data itself is biased or leads to biased interpretations, it can feed back into the AI training data, perpetuating or even amplifying the original biases. For example, predictive policing algorithms, when trained on historical biased crime data, can lead to over-policing in minority neighborhoods. This over-policing then generates more crime data from those areas, which subsequently reinforces the initial bias in the AI system. This interconnectedness underscores that ethical AI governance cannot be siloed. A solution for one ethical problem, such as mitigating bias, might inadvertently create or exacerbate another, such as privacy violations or new accountability gaps. This complex interplay necessitates a comprehensive, systems-thinking approach to AI ethics.

### **III. Bias and Discrimination in AI Systems**

AI bias refers to the systematic and unfair skewing of outcomes produced by artificial intelligence systems, often reflecting or amplifying existing societal prejudices.<sup>8</sup> This can lead to discriminatory results even when the AI system is not explicitly programmed to discriminate.<sup>10</sup>

#### **Sources of Bias**

Bias in AI systems can originate from several points within the AI lifecycle:

- **Data Bias:** This is widely considered the most prevalent source. Biases present in the historical data used to train AI models can lead to skewed outputs. If the training data predominantly represents certain demographics or contains discriminatory patterns, the AI will learn and perpetuate these imbalances in its predictions or decisions.<sup>3</sup> It is important to acknowledge that data collection is never neutral; it inherently reflects existing societal narratives and power structures.<sup>3</sup>
- **Algorithmic Bias:** This occurs when the design and parameters of algorithms inadvertently introduce bias, even if the underlying data were unbiased. The way algorithms process and prioritize certain features over others can result in discriminatory outcomes.<sup>8</sup>
- **Human Decision Bias (Cognitive Bias/User Bias):** Human prejudices and cognitive biases can seep into AI systems through subjective decisions made during data labeling, model development, and user interactions, whether consciously or unconsciously.<sup>8</sup> Developers' subjective choices in algorithm design and the trade-offs they make, such as prioritizing accuracy over fairness, also shape the system's outcomes.<sup>14</sup>
- **Generative AI Bias:** Generative AI models, used for creating text, images, or videos, can produce biased or inappropriate content based on the biases present in their vast training datasets. These models may reinforce stereotypes or generate outputs that marginalize certain groups or viewpoints.<sup>8</sup>

## Manifestations and Case Studies

AI bias manifests across various critical domains, leading to tangible discriminatory impacts:

- **Hiring and Recruitment (Amazon's tool):** Amazon's experimental AI recruiting tool serves as a prominent example. It was ultimately discontinued after it systematically discriminated against women applying for technical jobs. The algorithm, trained on ten years of predominantly male resumes, learned to downgrade CVs that included terms like "women's" or came from women's colleges, effectively perpetuating existing gender biases prevalent in the tech industry.<sup>3</sup> This case illustrates how AI can inadvertently "launder" human bias through software, making it appear objective while embedding historical prejudices.<sup>15</sup> Such practices have legal implications under Title VII, which prohibits employment discrimination.<sup>15</sup>
- **Criminal Justice (COMPAS algorithm):** The Correctional Offender Management Profiling for Alternative Sanctions (COMPAS) algorithm, widely used for recidivism

risk assessment in criminal sentencing, has been shown to exhibit racial bias. An analysis by ProPublica found that Black defendants were nearly twice as likely to be incorrectly flagged as high-risk compared to white defendants who did not re-offend, while white defendants were mistakenly labeled low-risk almost twice as often as Black re-offenders.<sup>8</sup> This case ignited significant debate regarding the definition of "fairness" in AI. The algorithm's creator argued for "accuracy equity," claiming the system predicted recidivism at similar rates for both groups. However, ProPublica focused on "error rate balance," highlighting the unequal rates of misclassification (false positives and false negatives) across racial groups.<sup>19</sup> This fundamental disagreement over what constitutes "fairness" reveals that AI bias is not solely a technical problem but a profound socio-ethical challenge. An algorithm can indeed be considered "fair" by one definition (e.g., predictive parity) while simultaneously being "unfair" by another (e.g., error rate balance), making universal mitigation strategies complex. The proprietary, "black box" nature of COMPAS further complicates scrutiny, making it difficult to verify its methodology, identify the sources of bias, or correct individual errors.<sup>20</sup> This opacity directly impedes accountability and effective bias mitigation, creating a significant hurdle for legal challenges and eroding public trust. If an individual is wrongfully impacted by a biased AI system, the lack of transparency makes it challenging to determine who is responsible or how the mistake occurred<sup>22</sup>, directly contributing to a "responsibility gap".<sup>23</sup>

- **Other Areas:** AI bias is not confined to hiring and criminal justice. It also manifests in credit scoring and lending, where algorithms may disadvantage certain socioeconomic or racial groups.<sup>8</sup> In healthcare, AI can introduce biases in diagnoses and treatment recommendations if trained on unrepresentative data.<sup>8</sup> Educational evaluation and admission algorithms can show bias against students from under-resourced backgrounds.<sup>8</sup> Predictive policing algorithms can lead to over-policing in minority neighborhoods.<sup>8</sup> Facial recognition systems often struggle with demographic accuracy, showing higher error rates for darker skin tones.<sup>3</sup> Voice recognition systems can exhibit bias against certain accents or dialects.<sup>8</sup> AI-based image generation systems may underrepresent or misrepresent certain racial or cultural groups.<sup>8</sup> Content recommendation algorithms can perpetuate echo chambers, reinforcing existing viewpoints.<sup>8</sup> Finally, insurance algorithms may unfairly determine premiums or eligibility based on biased factors like zip codes.<sup>8</sup>

## Mitigation Strategies: Technical, Policy, and Governance Solutions

Addressing AI bias requires a multi-faceted approach combining technical innovations with robust policy and governance frameworks:

- **Technical Strategies:** These include resampling methods (e.g., oversampling, undersampling) to balance imbalanced datasets, cost-sensitive learning, synthetic data generation (e.g., SMOTE, Generative Adversarial Networks - GANs, Variational Autoencoders - VAEs), and algorithmic modifications (e.g., ensemble methods, adjustments to loss functions).<sup>9</sup> Fairness-aware approaches, such as Equalized Odds and Demographic Parity, are also critical for mitigating bias, though their practical implementation often involves trade-offs with accuracy.<sup>26</sup>
- **Policy and Governance Solutions:** Robust data ethics frameworks, ethical AI governance structures, and continuous audits are crucial for responsible AI deployment.<sup>11</sup> This involves defining clear accountability mechanisms, fostering transparency, and implementing continuous model monitoring.<sup>11</sup> Establishing diverse AI ethics boards, incorporating Explainable AI (XAI) techniques to provide clear explanations for AI decisions, and maintaining comprehensive audit trails are key steps.<sup>11</sup> Bias awareness and critical thinking should be embedded from the conception phase of AI development, with diverse AI healthcare teams systematically reviewing for unintended consequences for specific demographic or socioeconomic groups.<sup>25</sup> Regulatory frameworks like the GDPR and CCPA also necessitate stringent governance practices to protect consumer rights and data privacy, which indirectly addresses bias by promoting responsible data handling.<sup>11</sup>

Table 1: Examples of AI Bias and Mitigation Strategies by Domain

Domain	Specific Manifestation of Bias	Case Study/Example	Source of Bias	Consequences/Impact	Mitigation Strategy (Technical/Policy/Governance)
Hiring	Gender bias in recruitment	Amazon Hiring Tool <sup>10</sup>	Data Bias (historical male-dominated data) <sup>10</sup>	Discriminatory outcomes against women <sup>10</sup>	Policy: Ethical AI governance, diverse AI teams. <sup>27</sup> Legal: Title VII implications. <sup>15</sup>

Criminal Justice	Racial bias in recidivism prediction	COMPAS algorithm <sup>18</sup>	Data Bias (historical crime data reflecting systemic biases) <sup>12</sup>	Black defendants misclassified as higher risk <sup>18</sup>	Technical: Fairness-aware algorithms (Equalized Odds). <sup>26</sup> Policy: Transparency, audit trails. <sup>11</sup>
Healthcare	Misdiagnosis for certain ethnic groups	AI trained on single ethnic group data <sup>8</sup>	Data Bias (unrepresentative training data) <sup>8</sup>	Inaccurate diagnoses, suboptimal care <sup>8</sup>	Technical: Diverse datasets, bias-proof features. <sup>25</sup> Policy: DEI principles in conception. <sup>27</sup>
Law Enforcement	Over-policing in minority neighborhoods	Predictive Policing Algorithms <sup>8</sup>	Data Bias (historical policing data) <sup>8</sup>	Exacerbated racial inequality, erosion of trust <sup>22</sup>	Policy: Transparency, accountability, explainability (XAI). <sup>11</sup>
Facial Recognition	Higher error rates for darker skin tones	Microsoft, IBM, Face++ algorithms <sup>3</sup>	Data Bias (training data lacking diversity) <sup>3</sup>	Misidentification, racial profiling <sup>3</sup>	Technical: Diverse datasets, improved algorithms. <sup>3</sup> Policy: Regulatory compliance, ethical guidelines. <sup>11</sup>
Generative AI	Reinforcement of stereotypes	Image generation systems <sup>8</sup>	Data Bias (biases in training data) <sup>8</sup>	Production of biased/inappropriate content <sup>8</sup>	Technical: Data augmentation, bias-aware

					algorithms. <sup>9</sup> Policy: Ethical guidelines for content generation. <sup>28</sup>
--	--	--	--	--	--

## IV. Autonomy and Accountability of AI Systems

### Defining Autonomy and Accountability in AI

Autonomy in AI refers to systems capable of learning, adapting, and making decisions without direct human input, often surpassing human capabilities in speed, precision, and data processing.<sup>23</sup> Accountability, in this context, calls for clear frameworks to define what is held accountable and to set up procedures for questioning and containing power in AI systems, specifically, the clear and unambiguous definition of responsibility lines for AI-generated decisions and actions.<sup>23</sup>

### The "Responsibility Gap": Challenges in Assigning Liability for AI Actions

The increasing autonomy of AI systems creates a significant "responsibility gap," where traditional legal doctrines, rooted in human-centric accountability (based on principles of intent, negligence, and foreseeability), are increasingly strained.<sup>23</sup> When autonomous systems, such as self-driving vehicles, cause harm, it becomes profoundly complex to determine who is legally liable: the software developer, the hardware manufacturer, the data provider, the end-user, or the AI system itself.<sup>23</sup> This challenge is exacerbated by the algorithmic opacity, often referred to as "black-box" systems, where the decision-making processes are opaque even to developers, and by the multi-party involvement typical in AI development and deployment.<sup>23</sup>

### Human-in-the-Loop vs. Human-on-the-Loop

The level of human involvement in AI systems is crucial for understanding accountability:

- **Human-in-the-Loop (HITL):** These systems represent a collaborative blend of AI



and human expertise, necessitating direct human intervention in decision-making processes. AI serves as an assistant, providing recommendations or executing tasks under human supervision. The goal is to combine AI's analytical capabilities with human judgment, especially in fields where the cost of errors is high, such as AI assisting radiologists in diagnosing medical images.<sup>34</sup> In HITL setups, AI systems drive the inference and provide suggestions, but humans ultimately intervene to provide corrections and supervision, retaining control of the full system.<sup>35</sup>

- **Human-on-the-Loop (HOTL) / AI-in-the-Loop (AI2L):** In these systems, AI is primarily in charge of decision-making, and human inputs are used to "guide" the model to a better optimum, making the process more efficient and possibly more effective. The human is not the ultimate decision-maker; rather, they oversee the system, with the full system existing independently of the AI's presence for its core function.<sup>35</sup> This distinction is critical for assigning decision-making authority and control.

The distinction between Human-in-the-Loop (HITL) and Human-on-the-Loop (HOTL) systems is paramount not only for technical design but also for legal and ethical accountability. A misunderstanding or mislabeling of a system's true level of autonomy can lead to inappropriate regulatory approaches and an unaddressed responsibility gap. When policymakers attempt to address concerns about algorithmic incapacities by simply "inserting a human into a decision-making process," they risk falling into what has been termed the "MABA-MABA trap" if they do not truly understand the *nature* and *degree* of human control. If a system is *de facto* AI-in-the-Loop (AI2L), where the AI makes the ultimate decisions, but is *labeled* as Human-in-the-Loop, it creates a false sense of human oversight. This can leave victims of AI-caused harm without clear recourse, as the legal framework might incorrectly assume human culpability where the machine was the primary decision-maker. This highlights that regulatory frameworks must precisely define the nature and degree of human involvement required for different risk levels of AI systems, rather than merely mandating a "human in the loop" without specifying the human's actual authority. It also points to the need for legal systems to evolve beyond traditional human-centric liability models to effectively address the complexities of autonomous AI.<sup>23</sup>

## Legal and Ethical Frameworks for Accountability

The legal landscape is evolving to address the complexities of AI accountability:

- **Evolving Legal Landscape:** Traditional tort law struggles to assign liability for



AI's unforeseen actions due to machine learning processes and the involvement of multiple parties (developers, manufacturers, data providers, users).<sup>23</sup> Legal scholars are actively debating solutions, including proposals for granting AI systems limited legal personhood (though critics argue this could allow companies to offload responsibility) or applying existing strict liability or vicarious liability doctrines.<sup>23</sup>

- **EU AI Act:** The European Commission's AI Act represents the most comprehensive regulatory attempt to date. It categorizes AI systems into risk tiers (unacceptable, high, limited, minimal) and assigns specific legal obligations based on these levels. High-risk systems, such as those used in critical infrastructure, employment, or law enforcement, are subject to stringent requirements, including mandatory transparency, human oversight, and conformity assessments throughout their lifecycle.<sup>23</sup> The Act also bans certain AI practices deemed to pose unacceptable risks, such as social scoring and real-time remote biometric identification in public spaces, with only very narrow exceptions for serious crimes and court approval.<sup>28</sup>
- **EU Liability Framework:** Complementing the AI Act, the EU proposes modernizing its liability framework to specifically address damages caused by AI systems. This includes introducing a rebuttable "presumption of causality" to ease the burden of proof for victims, making it easier for them to establish that damage was caused by an AI system.<sup>36</sup> Furthermore, national courts would be empowered to order the disclosure of evidence about high-risk AI systems suspected of having caused damage.<sup>36</sup> Operators of high-risk AI systems could be held liable for harm to life, health, or physical integrity of a natural person, damage to property, or significant immaterial harm resulting in verifiable economic loss.<sup>36</sup>

There is a fundamental tension between the need for legal certainty in AI accountability and the inherent "black box" nature and rapid evolution of AI systems. This makes static regulatory frameworks difficult to implement and enforce effectively. Many AI systems, especially proprietary ones, operate as "black boxes," obscuring their internal decision-making processes.<sup>12</sup> This opacity directly hinders the ability to assess bias and assign responsibility.<sup>12</sup> Traditional legal doctrines rely on proving fault, intent, or a clear causal link, which becomes exceedingly difficult when the inner workings of an AI system are inscrutable.<sup>23</sup> This opacity makes it challenging for victims to identify the liable party and prove the requirements for a successful liability claim.<sup>33</sup> While regulatory responses like the EU AI Act mandate transparency and human oversight for high-risk systems, and introduce a "presumption of causality" for damages, the rapid pace of AI innovation means that legal frameworks often struggle

to keep up.<sup>12</sup> This necessitates that legal frameworks be adaptive and incorporate mechanisms for continuous monitoring and auditing of AI systems.<sup>11</sup> Moreover, it suggests a potential shift in legal thinking from solely fault-based liability to strict liability for high-risk AI, acknowledging the inherent difficulty of proving intent or negligence in complex AI systems.<sup>23</sup>

- **Case Studies (Autonomous Vehicles):** Accidents involving self-driving cars vividly illustrate the complexities of AI accountability. While traditional car accidents typically attribute fault to the human driver, AI-driven accidents blur this line, raising questions about who is responsible when the machine makes decisions.<sup>31</sup> Incidents such as a self-driving taxi obstructing an ambulance, leading to a patient's death<sup>38</sup>, highlight the urgent need for clear liability rules. Legal solutions are being actively explored, including models where developers, manufacturers, and potentially the AI systems themselves are all held accountable.<sup>31</sup> Both state and federal legislative efforts, such as the SELF DRIVE Act in the US, are emerging to regulate autonomous vehicles, aiming to foster innovation while ensuring safety and reliability.<sup>38</sup>

## V. Job Displacement and the Future of Work

### The Threat of AI-Driven Automation: Scale and Scope of Potential Job Losses

Artificial intelligence-driven automation holds the potential for significant job displacement across various industries. Reports from investment banks like Goldman Sachs estimate that AI could replace the equivalent of 300 million full-time jobs globally, affecting approximately a quarter of all work tasks in the US and Europe.<sup>39</sup> McKinsey Global Institute analysts project that by 2030, at least 14% of employees globally could need to change their careers due to digitization, robotics, and AI advancements, with up to 30% of hours currently worked in the US potentially being automated.<sup>40</sup> Some estimates suggest that between 25% and 50% of jobs face significant AI exposure, impacting 40 to 80 million Americans in the coming years.<sup>42</sup> This prospect of widespread disruption is viewed as a serious and looming threat that demands immediate focus and attention.<sup>40</sup>

### Jobs Most and Least Susceptible to Automation

The impact of AI on the workforce is not uniform, with certain job categories being more vulnerable than others:

- **Most Susceptible:** Jobs involving repetitive, manual, or routine cognitive tasks that do not require high emotional or social intelligence are highly susceptible to automation. Examples include customer service representatives, where AI can provide automated responses to frequently asked questions.<sup>39</sup> Receptionists are increasingly being replaced by robots or AI-managed calls.<sup>39</sup> Accountants and bookkeepers are seeing their roles automated by AI-powered bookkeeping services that offer efficiency and cost savings.<sup>39</sup> Salespeople are affected as advertising shifts to web and social media, leveraging AI for targeted marketing.<sup>39</sup> The fields of data analysis and research are already implementing AI to streamline processes and identify new data without human assistance.<sup>39</sup> Warehouse work, insurance underwriting, and retail (e.g., self-checkout stations) are also areas experiencing significant automation.<sup>39</sup>
- **Least Susceptible:** Jobs requiring complex negotiation, strategy, emotional intelligence, leadership, creativity, and extensive personal experience are less likely to be fully replaced by AI. Teachers, for instance, often serve as a reference point and inspiration, making a fully digital teaching experience almost impossible in the foreseeable future.<sup>39</sup> Lawyers and judges rely heavily on negotiation, strategy, and case analysis, with a significant human factor involved in navigating complex legal systems and making critical decisions.<sup>39</sup> Directors, managers, and CEOs perform leadership roles that involve managing teams, sharing company missions, and fostering investor confidence, which are not easily codified into algorithms.<sup>39</sup> Similarly, HR managers perform vital tasks beyond hiring, such as maintaining staff motivation and managing discontent, requiring uniquely human skills.<sup>39</sup>

A significant observation is that AI's impact is shifting from primarily automating manual labor, characteristic of past industrial revolutions, to automating "mental" or "cognitive" labor. This represents a qualitative shift in technological disruption. Historically, technological change has always replaced old jobs with new ones.<sup>40</sup> However, the current wave of AI-driven automation is fundamentally different; it is not merely automating physical or repetitive tasks but increasingly automating cognitive labor, decision-making, and even creativity.<sup>43</sup> This is evident in the susceptibility of roles like accountants, research analysts, and insurance underwriters to automation.<sup>39</sup> This prospect of widespread automation of "mental" or "cognitive" labor could begin a process of "proletarianizing" portions of the "professional-managerial class".<sup>40</sup> This

constitutes a profound societal restructuring, as these classes were historically less vulnerable to technological displacement. This phenomenon challenges the traditional reassuring narrative that new jobs will always emerge to fully compensate for old ones, at least in terms of a smooth transition. The skills required for these new jobs—such as AI development, oversight, and integration—are often high-skilled<sup>41</sup>, thereby creating a growing divide between high-skilled and low-skilled workers.<sup>41</sup> This suggests a more profound societal transformation than previous industrial revolutions, with the potential for increased economic inequality and social disruption.<sup>40</sup>

### Job Creation and New Opportunities

Despite the concerns about job displacement, AI also creates new jobs and transforms existing roles.<sup>41</sup> The World Economic Forum predicts that AI and automation could contribute to the creation of 69 million new jobs worldwide by 2028.<sup>45</sup> Historically, new job titles that did not exist decades prior have emerged with new technologies, such as drone operators.<sup>40</sup> Firms that invest more in AI have shown associated growth in employment and heightened innovation, particularly in product development.<sup>44</sup> New roles emerging as a result of AI include:

- **AI Trainers and Teachers:** Individuals responsible for training and teaching AI systems.<sup>46</sup> This includes engineers and scientists developing large language models (LLMs), as well as electrical engineers for customized microchips and systems administrators for server infrastructure.<sup>47</sup>
- **AI Explainers:** These professionals design interfaces that enable the public to interact with AI, acting as "user experience designers" for LLMs. They may develop personalized AI assistants, tutors, or coaches.<sup>47</sup>
- **AI Sustainers:** This category ensures AI systems are used optimally. It includes content creators (e.g., prompt engineers who write text prompts for LLMs), data curators (responsible for ensuring high-quality training data), and ethics and governance specialists (who test systems for bias and ensure ethical use).<sup>47</sup>

### Policy Responses to Economic Disruption

Addressing the economic disruption caused by AI requires comprehensive policy responses:

- **Universal Basic Income (UBI):**
  - **Debate and Rationale:** Universal Basic Income is frequently discussed as a

potential solution to address wage inequality, job insecurity, and widespread job losses stemming from AI and automation.<sup>48</sup> It proposes a guaranteed income for all citizens, irrespective of their income, employment status, or other factors.<sup>49</sup>

- **Pilot Results and Feasibility:** Over 160 UBI tests and pilot programs have been conducted globally, generally yielding positive effects in alleviating poverty and improving health and education outcomes, though the evidence regarding impacts on employment is less clear.<sup>49</sup> Some studies report minimal to no significant reduction in work participation, while others indicate modest declines, often concentrated among secondary earners or those pursuing education or caregiving responsibilities.<sup>50</sup> While UBI can alleviate immediate financial stress, it is not a singular panacea for deeper systemic issues such as healthcare access, job stability, or upward mobility.<sup>50</sup> One study suggests that AI capital profits could sustainably finance a UBI, but this depends on the scale of AI productivity gains and the public revenue share from these profits.<sup>52</sup>
- A critical observation is that while Universal Basic Income (UBI) is frequently proposed as a comprehensive solution to AI-induced job loss, research suggests it is not a "panacea" and its impact on employment is mixed. This highlights a complex interplay rather than a simple solution. Evidence indicates that while UBI can alleviate poverty and improve well-being<sup>49</sup>, its efficacy regarding employment quality or upward mobility is limited.<sup>51</sup> Some studies show minimal work reduction, while others report modest declines in labor supply, often among specific groups like caregivers or students.<sup>50</sup> This implies that UBI, while valuable as a social safety net, cannot be the sole policy response to AI's impact on the workforce. It needs to be integrated as part of a broader, multi-faceted policy toolkit that also includes significant investment in education, skills development, and active labor market policies.<sup>50</sup> The ongoing debate also underscores the importance of UBI's specific design, including the level of income provided and its funding mechanisms, in determining its overall effectiveness.<sup>50</sup>
- **Workforce Retraining and Upskilling Initiatives:**
  - **Challenges:** The current education system is "dangerously out of sync" with the speed and scale of AI advancements, characterized by outdated curricula, delayed approvals, and institutional inertia.<sup>42</sup> Curriculum changes can take 18 months or more, and faculty incentives often reward theoretical knowledge over cutting-edge practical application.<sup>42</sup> Furthermore, many retraining courses are expensive, often out of reach for the very workers most likely to face displacement, and existing workforce systems are not designed for the

massive scale of retraining needed (40-80 million people).<sup>42</sup>

- **Approaches:** Policy recommendations include encouraging companies to retrain workers, making health benefits portable, reducing vesting requirements for retirement benefits, loosening job licensing requirements (especially for non-health/safety related fields), and creating worker retraining accounts that function similarly to retirement accounts.<sup>53</sup> Policymakers must also ensure universal access to high-speed internet to facilitate participation in retraining programs.<sup>53</sup> New approaches involve experimenting with hybrid learning formats, targeted micro-credential courses, and AI-enabled personalized learning platforms that can adapt to new market demands in real-time.<sup>42</sup> Developing better labor market projections and collecting robust evidence on what works in retraining programs, particularly for technology-displaced workers, are also crucial.<sup>54</sup>
- A critical systemic challenge is the growing mismatch between the rapid pace of AI development and the slow-moving traditional education and policy systems. This disparity could lead to a significant "unemployment crisis" if not addressed urgently. AI models iterate rapidly, with new advancements emerging frequently <sup>42</sup>, and there is an increasing sense that "time is running out" for effective action.<sup>55</sup> In contrast, the education system is burdened by "outdated curricula," "delayed approvals," and "institutional drag".<sup>42</sup> Curriculum changes require lengthy processes, and faculty incentives often prioritize theoretical knowledge over practical, cutting-edge applications.<sup>42</sup> This disconnect means that by the time workers are trained in current AI technologies, the technologies themselves may have evolved significantly, rendering the training partially obsolete.<sup>42</sup> This creates a critical bottleneck for workforce adaptation. The potential consequence is the "largest unemployment crisis since the Great Depression" if education and retraining systems fail to catch up.<sup>42</sup> This highlights the need for radical shifts in how education and workforce development are funded, structured, and delivered, moving towards more agile, tech-driven training models that can adapt in weeks, not years.<sup>42</sup> It also implies that traditional political cycles are too slow to respond to the accelerating pace of AI-driven change.<sup>42</sup>

Table 2: AI's Impact on Job Roles: Automation vs. Resilience

Job Category	Susceptibility to Automation	Reason for Susceptibility/Resilience	Specific AI Impact	New/Emerging Roles	Industry

Customer Service Representative	High	Repetitive queries, low emotional intelligence <sup>39</sup>	Automated responses, chatbots <sup>39</sup>	-	Service
Accountant/Bookkeeper	High	Routine data processing, formula application <sup>39</sup>	AI-powered bookkeeping, automated analysis <sup>39</sup>	-	Finance
Warehouse Worker	High	Repetitive physical tasks, logistics optimization <sup>39</sup>	Mechanized retrieval, automated sorting <sup>39</sup>	-	Logistics
Research & Analysis	High	Efficient data sorting, extrapolation, pattern identification <sup>39</sup>	AI-driven data analysis, automated research <sup>39</sup>	Data Curators <sup>47</sup>	Various
Teacher	Low	Requires inspiration, emotional intelligence, complex academic decisions <sup>39</sup>	AI as personalized tutors/assistants <sup>47</sup>	AI Explainers (tutors) <sup>47</sup>	Education
Lawyer/Judge	Low	Complex negotiation, strategy, human judgment in legal systems <sup>39</sup>	AI for document review, predictive analytics <sup>17</sup>	-	Legal
CEO/Manager	Low	Leadership, team motivation, strategic	AI for efficiency, data-driven	-	Various



		vision, investor confidence <sup>39</sup>	decisions <sup>45</sup>		
AI Trainer	N/A (New Role)	Expertise in developing/teaching AI systems <sup>46</sup>	Directly involved in AI development <sup>46</sup>	AI Trainer <sup>46</sup>	Tech
Prompt Engineer	N/A (New Role)	Crafting effective prompts for LLMs <sup>47</sup>	Enables optimal AI content generation <sup>47</sup>	Prompt Engineer <sup>47</sup>	Tech/Creative
AI Ethicist/Governance Specialist	N/A (New Role)	Ensuring ethical use, mitigating bias, compliance <sup>47</sup>	Oversees responsible AI deployment <sup>47</sup>	AI Safety Officer, Ethicist <sup>47</sup>	Tech/Policy

VI. Privacy and Surveillance in the Age of AI

AI-Powered Surveillance Technologies: Facial Recognition, Predictive Policing, Data Mining

AI systems are increasingly integrated into surveillance infrastructures, enabling real-time interpretation of footage, tracking movements, scanning individuals for weapons, analyzing behavior, and monitoring access to buildings.<sup>56</sup> This includes widely deployed technologies such as facial recognition in public spaces, predictive policing algorithms, and social media monitoring.<sup>22</sup> AI's unparalleled ability to process and analyze vast amounts of personal data allows it to deduce sensitive information about individuals, including sexual orientation, political convictions, or health concerns, which could then be exploited.<sup>57</sup>

Concerns and Misuse Cases

The proliferation of AI in surveillance raises significant ethical and societal concerns:

- **Privacy Violations:** AI-driven surveillance can lead to a profound erosion of privacy and civil liberties through constant tracking and profiling of individuals without their knowledge or consent.<sup>22</sup> Examples include facial recognition systems misidentifying people of color and women at higher rates<sup>22</sup>, leading to concerns about racial profiling. The use of AI on university campuses, for instance, has raised concerns among students and activists about potential surveillance of their movements and expressions.<sup>56</sup>
- **Data Breaches and Exploitation:** AI's reliance on vast amounts of sensitive personal data—including financial data, health records, behavioral patterns, and biometric information—significantly increases the risk of data breaches and exploitation for unintended purposes.<sup>58</sup> A notable case involved ProctorU, an online exam proctoring service, which experienced a data breach that leaked user records for approximately 444,000 students.<sup>59</sup> Another instance involved alleged HIPAA violations by Google/DeepMind for accessing countless patient medical files for AI data-mining without proper consent.<sup>60</sup>
- **Constant Monitoring and Trust Erosion:** Continuous monitoring by AI systems can create a pervasive sense of surveillance, potentially leading individuals to self-censor their thoughts and behaviors and eroding trust in the institutions employing such technologies.<sup>59</sup>
- **Deepfakes:** Generative AI has enabled the creation of sophisticated synthetic media, known as deepfakes, including videos, audio, images, and text. These can convincingly distort reality, undermine public confidence in media, and be weaponized for malicious purposes, such as creating false narratives or inflammatory speeches attributed to public figures.<sup>7</sup>
- **Enforcement Actions:** Regulatory bodies have begun taking action against AI misuse. The US Federal Trade Commission (FTC) has acted against companies like Rite Aid for using facial recognition technology that falsely tagged consumers, particularly women and people of color, as shoplifters.<sup>62</sup> Amazon's voice assistant service, Alexa, faced a complaint over its default settings to retain users' voice recordings indefinitely without clear consent.<sup>62</sup> Clearview AI faced multiple lawsuits under the Biometric Information Privacy Act (BIPA) and the California Consumer Privacy Act (CCPA) for collecting biometric data without informing consumers.<sup>60</sup> Even municipalities, such as the city of Trento in Italy, have been fined for AI privacy violations.<sup>60</sup>

## Regulatory Frameworks: GDPR, CCPA, and their Provisions for AI

Key regulatory frameworks are emerging to address AI privacy concerns:

- **GDPR (General Data Protection Regulation):** The GDPR sets stringent requirements for data protection. It mandates explicit opt-in consent for data collection and processing, emphasizes transparency regarding how data is processed, kept, and used, and grants individuals several rights. These rights include the right to know about data processing, access personal data, correct inaccurate information, delete data, give prior permission, revoke consent, and the right not to be subject to decisions based solely on automated processing.<sup>63</sup> The GDPR broadly applies to any fully automated decision-making process that produces legal or similarly significant effects, such as affecting employment or creditworthiness, and treats sensitive data with heightened protection.<sup>64</sup>
- **CCPA (California Consumer Privacy Act):** In contrast to GDPR's opt-in model, the CCPA generally allows an opt-out consent model for data collection.<sup>63</sup> It grants consumers rights to know what personal data companies own, to delete collected data, and to opt-out of the sale of personal data.<sup>63</sup> Recent amendments aim to regulate Automated Decision-Making Technologies (ADMTs), requiring pre-use notices for impactful decisions (e.g., hiring, loan approvals) and updates to privacy policies to inform consumers of their opt-out rights.<sup>65</sup>
- **Common Principles:** Both regulations emphasize principles like "privacy-by-design," which integrates privacy considerations into the design of AI systems from the outset, data minimization (only collecting the bare minimum of data needed), purpose limitation, and overall transparency in data handling.<sup>63</sup> They aim to elevate data privacy from a mere compliance task to a core ethical principle, fostering the development of AI systems that are secure, transparent, and future-ready.<sup>63</sup>

### Privacy-Preserving AI Techniques: Federated Learning, Differential Privacy

To mitigate privacy risks while leveraging AI's capabilities, several advanced techniques are being developed:

- **Federated Learning (FL):** This is a distributed machine learning paradigm designed to preserve user privacy by allowing AI models to be trained on decentralized data sources, such as user devices, without centralizing or sharing the raw data with a central server. This approach enables individuals to maintain control over their personal data while still benefiting from AI advancements.<sup>66</sup>
- **Differential Privacy (DP):** A rigorous mathematical framework that bolsters privacy guarantees by introducing calibrated noise to model updates or data

queries. It ensures that information about a single individual's data is protected, making it difficult to infer specific details about any one person from the aggregated results, even when information from the dataset is publicly available.<sup>66</sup>

- **Trade-offs and Challenges:** While these privacy-preserving methods are crucial for mitigating risks, challenges persist in balancing privacy, communication efficiency, and model accuracy.<sup>66</sup> Differential Privacy, for example, can introduce computational overhead, leading to increased memory usage, computational variation across devices, and higher battery consumption, especially when strict privacy budgets (requiring more noise) are enforced.<sup>69</sup> This can further impact the model's accuracy.<sup>69</sup> Other secure computational techniques, such as Multi-Party Computation (MPC) and Homomorphic Encryption (HE), also offer strong privacy guarantees but introduce significant computational complexity and overhead, often making them impractical for resource-constrained edge devices or large-scale deployments.<sup>66</sup>

There is an inherent tension between maximizing AI's utility, which often requires vast amounts of data for training and performance optimization, and ensuring individual data privacy. This necessitates a trade-off that current privacy-preserving techniques are still grappling with. AI systems improve significantly with access to extensive and diverse datasets.<sup>57</sup> However, this extensive data collection inherently raises significant privacy concerns.<sup>22</sup> While privacy-preserving techniques like Federated Learning (FL) and Differential Privacy (DP) are designed to address this by allowing models to be trained without centralizing raw data or by adding noise to protect individual data <sup>66</sup>, these methods are not without their costs. As noted, challenges persist in balancing "privacy, communication efficiency, and model accuracy".<sup>66</sup> Differential Privacy, for instance, can lead to increased memory usage, computational cost, and battery consumption, particularly when stricter privacy guarantees are sought.<sup>69</sup> This implies that achieving robust privacy in AI often comes at a cost to model performance or computational efficiency. Policymakers and developers face difficult decisions about acceptable trade-offs based on the application's sensitivity and the potential impact on individuals. This underscores that privacy is not a simple compliance checkbox but a complex engineering and policy challenge requiring ongoing research and development to optimize these inherent trade-offs.<sup>66</sup>

**Table 3: Key Privacy Provisions for AI under GDPR and CCPA**

Aspect	GDPR Provision/Approach	CCPA Provision/Approach	Key Differences/Similarities	Implications for AI Systems

<b>Consent Model</b>	<b>Opt-in:</b> Informed consent mandatory before data collection <sup>63</sup>	<b>Opt-out:</b> Consumers can choose not to have data used in ADMTs under certain circumstances <sup>63</sup>	GDPR is stricter (explicit opt-in); CCPA offers more flexibility (opt-out) <sup>63</sup>	AI systems must integrate robust consent mechanisms, varying by jurisdiction; GDPR requires higher bar for data use <sup>63</sup>
<b>Rights Granted to Individuals</b>	Right to know, access, correct, delete data; right to revoke consent; right not to be subject to solely automated decisions <sup>63</sup>	Right to know, delete, opt-out of data sales; right to non-discrimination <sup>63</sup>	GDPR includes specific right against automated decisions; CCPA focuses on data control <sup>63</sup>	AI developers must build systems that respect these rights, enabling data access, deletion, and challenging automated decisions <sup>63</sup>
<b>Scope of Automated Decision-Making (ADM)</b>	Applies to any fully automated decision-making process with legal or similarly significant effects <sup>64</sup>	Focuses on ADMTs that replace human judgment or are key factors in significant decisions (e.g., hiring, loans) <sup>65</sup>	GDPR casts a broader net; CCPA has more carve-outs and focuses on "impactful" decisions <sup>65</sup>	AI systems in high-stakes contexts require pre-use notices, explanations, and potentially human review, especially in the EU <sup>64</sup>
<b>Sensitive Data Treatment</b>	Heightened protection; requires stronger justifications for use <sup>64</sup>	Does not clearly adopt a layered approach for sensitive data <sup>65</sup>	GDPR provides stronger safeguards for biometric, health, etc. data <sup>65</sup>	AI models processing sensitive data face higher compliance burdens under GDPR, necessitating stricter privacy-by-design <sup>65</sup>
<b>Transparency &amp;</b>	Requires clear information on data processing	Requires pre-use notices for ADMTs; right	Both emphasize transparency; GDPR provides	AI systems need to be designed for explainability

<b>Explainability</b>	purpose; uncertainties exist on individual explanation for ADM <sup>63</sup>	to explanation/appeal is conditional <sup>65</sup>	unconditional right to human review/explanation for ADM <sup>65</sup>	(XAI) to meet regulatory demands and build trust, especially for impactful decisions <sup>11</sup>
-----------------------	--	--	---	--

## VII. AI and Weapons Development

### Autonomous Weapons Systems (AWS): Definition and Ethical Concerns

Autonomous Weapons Systems (AWS) are defined as weapons designed to independently select and engage targets based on sensor processing, without the need for manual human intervention.<sup>71</sup> Colloquially known as "killer robots," these systems have become a central focus of international ethical debates.<sup>73</sup> The ethical concerns surrounding AWS are profound and multifaceted. They include the potential for an increase in wrongs and crimes during military operations, the erosion of moral responsibility among human operators, and the creation of a "responsibility gap" for harms caused by these systems.<sup>24</sup> AWS raise fundamental questions about the morality of delegating life-and-death decisions to machines that inherently lack human judgment, the capacity to understand or respect the true value of human life, or the ability for moral agency.<sup>30</sup>

### Implications for Warfare: Erosion of Moral Responsibility, Accountability Gap, Human Control

The proliferation of AI in warfare introduces unprecedented challenges, fundamentally reshaping humanity's relationship with violence.<sup>76</sup> A key concern is the potential for automation bias and technological mediation to weaken moral agency among operators of AI-enabled targeting systems, thereby diminishing their capacity for ethical decision-making in critical situations.<sup>76</sup> Autonomous weapons systems, by their nature, would be unable to identify subtle cues of human behavior, interpret the necessity of an attack, weigh proportionality, or communicate effectively to defuse a situation and ensure lethal force is a last option.<sup>71</sup> This raises serious concerns about their potential to violate international human rights law.<sup>71</sup>

The "responsibility gap" is a particularly pressing issue, as it remains unclear who could be held legally responsible if AWS violate international humanitarian law or human rights law.<sup>74</sup> There are significant legal challenges to holding individual operators criminally liable for the unpredictable actions of a machine they cannot fully understand, and equally substantial hurdles to finding programmers and developers responsible under civil law.<sup>74</sup> This lack of clear accountability undermines the fundamental principles of justice and trust in the use of force.

### **The Debate on Meaningful Human Control (MHC)**

In response to these concerns, the principle of "meaningful human control" (MHC) has emerged as a central concept in the legal-political debate surrounding AWS.<sup>24</sup> This principle asserts that humans, not AI, should ultimately remain in control of, and thus morally responsible for, relevant decisions about lethal military operations.<sup>24</sup> However, a detailed philosophical and technical theory of what MHC exactly means, and how it can be implemented in practice, is still lacking for policymakers and technical designers.<sup>24</sup> Despite this definitional challenge, MHC is widely considered a central notion in the ethics of robotics and AI, particularly for ensuring that human values and moral agency are preserved in the deployment of autonomous systems.<sup>24</sup>

### **International Law and Treaty Negotiations: Status and Challenges**

The international community has been grappling with the regulation of lethal autonomous weapons systems for over a decade.

- **Current Status:** Discussions on LAWS have been ongoing since 2014 under the Convention on Certain Conventional Weapons (CCW) in Geneva.<sup>55</sup> In December 2024, the United Nations General Assembly adopted a resolution on LAWS with overwhelming support (166 votes in favor), which mentioned a potential two-tiered approach: prohibiting some lethal autonomous weapon systems while regulating others under international law.<sup>75</sup> More recently, in May 2025, a UN General Assembly meeting saw participation from 96 countries, where the UN Secretary-General and the President of the International Committee of the Red Cross (ICRC) reiterated their call for the conclusion of a legally binding instrument banning these "politically unacceptable, morally repugnant" weapons by 2026.<sup>55</sup> Over 120 countries support calls to negotiate a treaty that prohibits AWS operating without meaningful human control or those that target people.<sup>74</sup>
- **Challenges and Sticking Points:** Despite growing momentum, substantive



progress in the CCW has been hindered by its reliance on a consensus decision-making model.<sup>74</sup> This approach allows a small number of major military powers investing in autonomous weapons systems—most notably India, Israel, Russia, and the United States—to repeatedly block proposals for a legally binding instrument.<sup>74</sup> Key sticking points remain, including the lack of a universally accepted definition of LAWS and a detailed understanding of what "meaningful human control" truly entails in practice.<sup>55</sup> Critically, talks are currently consultations only and "are not yet negotiating" a legally binding agreement, despite the increasing urgency.<sup>55</sup> The rapid evolution of AI and the decreasing costs of developing autonomous systems mean that "time is running out" to establish effective regulations before proliferation becomes widespread.<sup>55</sup>

There is a critical and growing disconnect between the rapid technological advancement and proliferation of autonomous weapons systems and the slow, consensus-driven pace of international legal and ethical regulation. This creates a dangerous "governance gap." AI advances are astonishingly rapid<sup>6</sup>, spurring the quick development of autonomous weapons systems.<sup>74</sup> These systems are also becoming less expensive to develop, raising concerns about their proliferation among both state and non-state actors.<sup>55</sup> In stark contrast, international discussions on LAWS have been ongoing for over a decade<sup>55</sup> but have yielded "no substantive outcome" due to the Convention on Conventional Weapons' (CCW) consensus model.<sup>74</sup> This allows a single country to block proposals, meaning that talks are still "consultations only" and "not yet negotiating" a legally binding agreement.<sup>55</sup> This disparity in pace means that the technology is outpacing the development of legal and ethical frameworks<sup>22</sup>, leading to a substantial governance gap. The inherent "uncertainty of future developments"<sup>72</sup> further complicates regulation, yet the "looming threat" demands immediate focus.<sup>40</sup> This situation creates a high-stakes environment where potentially "morally repugnant" weapons<sup>55</sup> could become widespread before adequate international norms or prohibitions are in place. The current lack of transparency in weapons development<sup>72</sup> exacerbates this problem. The urgent declaration that "time is running out"<sup>55</sup> is a critical message for policymakers, highlighting the need for more agile and decisive international action beyond the limitations of traditional consensus-based forums.

## **VIII. Overarching Ethical Principles and Stakeholder Responsibilities**

## Core Ethical Principles for AI

A broad consensus is emerging around several core ethical principles intended to guide the responsible development and deployment of AI systems:

- **Transparency/Explicability:** This principle emphasizes the need to understand how an AI system arrives at its decisions, especially when those decisions have significant impacts on individuals or society.<sup>2</sup> It involves providing clear explanations for specific AI decisions and thoroughly documenting algorithms and data sources.<sup>11</sup>
- **Justice and Fairness:** This aims to prevent discrimination and ensure equitable treatment for all users. It often involves developing metrics to detect and mitigate bias in training data and model outputs.<sup>2</sup>
- **Accountability/Responsibility:** This principle assigns responsibility when something goes wrong with an AI system. It entails establishing clear audit trails for AI actions and holding developers and deployers accountable for discriminatory or harmful outcomes.<sup>2</sup>
- **Beneficence/Non-Maleficence:** Rooted in bioethics, this principle dictates that AI technology should promote human well-being, preserve human dignity, and avoid causing harm.<sup>3</sup> It extends to ensuring AI contributes to broader societal well-being and environmental sustainability.<sup>79</sup>
- **Privacy:** This involves protecting sensitive personal information used in AI processes, ensuring user control over their data, and adhering to established data protection regulations.<sup>2</sup>
- **Human Dignity and Autonomy:** This principle emphasizes that AI systems should respect individual autonomy and not undermine human agency or the unique value of human creative expression.<sup>3</sup>

A significant challenge lies in translating abstract ethical principles like "fairness" or "transparency" into concrete technical specifications and practical deployment guidelines. This process is inherently complex, often involving difficult trade-offs and varying interpretations across different contexts and cultures. For example, defining "fairness" objectively across diverse user groups and application contexts presents a considerable hurdle.<sup>77</sup> Similarly, effectively explaining the decision-making processes of complex "black box" deep learning models to ensure transparency remains a technical and communicative challenge.<sup>77</sup> The very definitions of "harm" or "fairness" can vary significantly across different societal norms, necessitating a decolonized approach to AI ethics that respects pluralism and avoids replicating historical biases

on a global scale.<sup>77</sup> This implies that ethical AI development is not a purely technical problem solvable by engineers alone. It requires continuous evaluation, adaptation, and interdisciplinary collaboration involving ethicists, social scientists, legal experts, and diverse stakeholders.<sup>14</sup> The absence of a unified ethical standard across global regions also complicates international AI governance and cross-border enforcement, highlighting the need for ongoing dialogue and harmonization efforts.<sup>2</sup>

Table 4: Foundational Ethical Principles for AI

Principle	Definition/Core Concept	Intermediate Application	Potential Challenge in Implementation	Relevant Snippet IDs
Fairness	Preventing discrimination and ensuring equitable treatment for all users <sup>11</sup>	Developing metrics to detect and mitigate bias in training data and model outputs <sup>11</sup>	Defining 'fairness' objectively across different contexts and user groups <sup>77</sup>	<sup>11</sup>
Transparency	Understanding how an AI system arrives at its decisions <sup>11</sup>	Providing clear explanations for specific AI decisions (e.g., loan denial) <sup>11</sup>	Explaining complex 'black box' deep learning models effectively <sup>12</sup>	<sup>11</sup>
Accountability	Assigning responsibility when something goes wrong <sup>11</sup>	Establishing audit trails for AI actions and assigning responsibility within organizations <sup>11</sup>	Identifying the responsible party in complex, multi-component AI systems <sup>23</sup>	<sup>11</sup>
Beneficence	Promoting well-being, preserving dignity, and sustaining the planet <sup>79</sup>	Designing AI for the common good and benefit of humanity <sup>79</sup>	Balancing diverse interpretations of "well-being" and "common good" across stakeholders <sup>79</sup>	<sup>61</sup>

<b>Privacy</b>	Protecting sensitive information used in AI processes <sup>63</sup>	Implementing formal data governance processes, using privacy-by-design techniques <sup>58</sup>	Balancing privacy with communication efficiency and model accuracy <sup>66</sup>	2
<b>Human Dignity</b>	Ensuring AI respects individual autonomy and human agency <sup>77</sup>	Avoiding cognitive behavioral manipulation or social scoring <sup>28</sup>	Defining and protecting human purpose and fulfillment in an AI-mediated world <sup>77</sup>	3

### Roles of Stakeholders: Governments, Corporations, Academia, Civil Society

Effective AI governance requires the active participation and collaboration of diverse stakeholder groups:

- Governments:** Play a crucial role in developing national strategies, guidelines, and enacting laws to regulate AI use, address bias, and protect privacy.<sup>37</sup> Examples include the EU AI Act, GDPR, and CCPA.<sup>37</sup> However, legislative processes are often slow, struggling to keep pace with the rapid advancements in AI.<sup>37</sup>
- Corporations:** Faced with mounting pressure from ethical challenges and potential risks, businesses are increasingly engaging in self-scrutiny and internal regulation. Many companies have defined and publicized their ethical principles, formed AI ethics boards, provided ethics training to employees, and formed industry alliances to promote ethical AI development.<sup>37</sup> They are recognizing their obligation to comply with external regulations and build trust with consumers and the broader public.<sup>37</sup>
- Academia:** Research institutions and think tanks actively study the ethical implications of AI, provide policy recommendations, and serve as independent checks on private sector activities. They provide the theoretical underpinnings necessary to understand AI's ethical dimensions.<sup>37</sup> However, academic research often struggles to keep pace with the rapid development of AI, creating a gap between theoretical insights and practical industry application.<sup>37</sup>
- Civil Society/NGOs:** Organizations such as Human Rights Watch and the Stop Killer Robots campaign play a vital role in advocating for human rights protection and pushing for specific prohibitions and regulations on AI, particularly in

high-stakes areas like autonomous weapons systems.<sup>55</sup> They are crucial in raising public awareness and pushing for higher ethical standards.

An emerging theme in AI governance is the recognition that no single stakeholder group can effectively govern AI ethics alone; a multi-stakeholder, interdisciplinary, and collaborative approach is not just beneficial but essential for comprehensive and adaptive governance. Governments, with their slow legislative processes, often lag behind technological advancements.<sup>37</sup> Academia, while providing critical theoretical foundations, struggles to keep pace with the rapid evolution of AI.<sup>37</sup> Corporations, despite increasing self-regulation, may not go beyond a shared understanding of ethical principles without external pressure.<sup>37</sup> The complexity and rapid evolution of AI, coupled with the interconnectedness of its ethical challenges, mean that a fragmented approach by individual stakeholders is insufficient. Each group brings unique expertise and perspectives—technical knowledge from industry, legal frameworks from government, and ethical/societal considerations from academia and civil society. This necessitates "interdisciplinary collaboration between policymakers, computer scientists, ethicists and social scientists"<sup>14</sup> and "collaborative platforms involving researchers, policymakers, and technologists".<sup>1</sup> The future of effective AI governance lies in fostering robust partnerships and shared responsibility across these diverse groups to create adaptive and comprehensive ethical frameworks.<sup>37</sup>

## IX. Conclusion and Recommendations

The rapid advancement of artificial intelligence presents both unprecedented opportunities for societal advancement and profound ethical challenges that demand urgent attention. As explored throughout this report, issues of bias, accountability, job displacement, privacy, and autonomous weapons are deeply intertwined and cannot be addressed in isolation. The complexities range from the fundamental disagreements on what constitutes "fairness" in AI outcomes to the profound legal ambiguities surrounding liability for autonomous systems. The rapid pace of AI development consistently outstrips the capacity of traditional regulatory and educational systems to adapt, creating dangerous governance gaps and potential societal disruptions.

To navigate this ethical frontier responsibly, a multi-pronged approach involving all stakeholders is imperative:

- **Policy & Regulation:**
  - Develop agile, adaptive regulatory frameworks that can keep pace with

technological advancements, potentially moving beyond slow, consensus-based models for high-stakes areas like autonomous weapons systems.

- Mandate transparency and explainability (XAI) for high-risk AI systems, along with continuous auditing mechanisms, to foster trust and enable effective oversight.
- Establish clear and robust liability frameworks that effectively address the "responsibility gap" for autonomous systems, ensuring accountability when harm occurs.
- Implement comprehensive data privacy laws with strong enforcement mechanisms, building upon principles like data minimization, purpose limitation, and user consent.

- **Industry & Development:**

- Prioritize "ethics-by-design" from the earliest conception phase of AI development, embedding Diversity, Equity, and Inclusion (DEI) principles and bias awareness throughout the entire AI lifecycle.
- Invest significantly in privacy-preserving AI techniques such as Federated Learning and Differential Privacy, alongside dedicated research to mitigate the inherent trade-offs with model accuracy and efficiency.
- Foster diverse and interdisciplinary AI development teams to bring varied perspectives and minimize unintentional biases.
- Implement robust internal governance structures, establish independent ethics boards, and provide ongoing ethical training for all employees involved in AI development and deployment.

- **Workforce & Society:**

- Invest substantially in accessible, adaptive workforce retraining and upskilling programs that are designed to meet the evolving skills gap created by AI. These programs must be agile and responsive to changing market demands.
- Explore and implement social safety nets, such as Universal Basic Income (UBI), as part of a broader, multi-faceted response to economic disruption, while acknowledging its limitations as a sole solution.
- Promote widespread public literacy and critical thinking regarding AI to empower individuals to understand and navigate its impacts.

- **Research & Collaboration:**

- Continue interdisciplinary research into defining and measuring fairness in AI, developing more effective bias mitigation techniques, and optimizing privacy-preserving methods to overcome current trade-offs.
- Foster strong and continuous collaboration between academia, industry, government, and civil society to share insights, develop best practices, and

collectively navigate the complex ethical landscape of AI.

## Future Outlook on Emerging Ethical Challenges

As AI capabilities continue to advance, several nascent and speculative ethical considerations are poised to gain increasing prominence, demanding proactive foresight rather than reactive policy-making:

- **Artificial General Intelligence (AGI) and Superintelligence:** The ethical implications of AI systems achieving human-level or superhuman intelligence are profound, raising fundamental questions about control, alignment with human values, and potential existential risks.<sup>3</sup>
- **AI Consciousness and Moral Status:** Philosophical and ethical debates are intensifying around the possibility of conscious AI systems. This encompasses their potential capacity for suffering, and the question of whether they deserve moral consideration or even rights, pushing the boundaries of traditional ethical frameworks.<sup>3</sup>
- **Societal Integration and Human-AI Coexistence:** Deeper questions will emerge about human identity, agency, creativity, and purpose in a world increasingly mediated by sophisticated artificial agents. This involves re-evaluating concepts of authorship, originality, and the unique value of human creative expression.<sup>3</sup>

The discussion of future challenges like Artificial General Intelligence (AGI) and AI consciousness highlights the critical need for proactive ethical foresight, rather than reactive policy-making, to address potential risks before they materialize. The report has consistently shown how current ethical challenges, such as bias and accountability, often arise because technological advancements outpace regulatory and societal adaptation.<sup>12</sup> If current, narrower AI applications already pose significant ethical dilemmas that society struggles to govern effectively, then the advent of AGI or conscious AI, with their far-reaching implications for the very nature of intelligence, consciousness, and human existence, demands an even greater level of anticipatory ethical planning. This implies that the ethical discourse and governance frameworks for AI must evolve from merely addressing

*current harms to actively anticipating and mitigating future, potentially catastrophic risks.* This requires dedicated investment in "AI safety and alignment" research<sup>3</sup> and the establishment of principles for responsible research and deployment, even if the possibility of AI consciousness remains theoretical.<sup>5</sup> The imperative shifts from "fixing problems" to "preventing existential risks"<sup>3</sup>, underscoring the urgent need for a



forward-looking and preventative approach to AI ethics.

## Works cited

1. Addressing Bias, Privacy, and Job Displacement in AI Integration - IJCRT, accessed on July 17, 2025, <https://www.ijcrt.org/papers/IJCRT2501262.pdf>
2. AI Ethics: Integrating Transparency, Fairness, and Privacy in AI Development, accessed on July 17, 2025, [https://www.researchgate.net/publication/388803359\\_AI\\_Ethics\\_Integrating\\_Transparency\\_Fairness\\_and\\_Privacy\\_in\\_AI\\_Development](https://www.researchgate.net/publication/388803359_AI_Ethics_Integrating_Transparency_Fairness_and_Privacy_in_AI_Development)
3. Ethics of artificial intelligence - Wikipedia, accessed on July 17, 2025, [https://en.wikipedia.org/wiki/Ethics\\_of\\_artificial\\_intelligence](https://en.wikipedia.org/wiki/Ethics_of_artificial_intelligence)
4. (PDF) Ethical Challenges in the Integration of AGI in Scientific Research - ResearchGate, accessed on July 17, 2025, [https://www.researchgate.net/publication/378831058\\_Ethical\\_Challenges\\_in\\_the\\_Integration\\_of\\_AGI\\_in\\_Scientific\\_Research](https://www.researchgate.net/publication/378831058_Ethical_Challenges_in_the_Integration_of_AGI_in_Scientific_Research)
5. Principles for Responsible AI Consciousness Research, accessed on July 17, 2025, <https://jair.org/index.php/jair/article/view/17310>
6. Principles for Responsible AI Consciousness Research - arXiv, accessed on July 17, 2025, <https://arxiv.org/pdf/2501.07290>
7. Debating the ethics of deepfakes, accessed on July 17, 2025, <https://www.orfonline.org/expert-speak/debating-the-ethics-of-deepfakes>
8. AI Bias - Artificial Intelligence in Education - LibGuides at Marian University, accessed on July 17, 2025, <https://libguides.marian.edu/c.php?g=1321167&p=10767259>
9. Fairness and Bias in Artificial Intelligence: A Brief Survey of Sources, Impacts, and Mitigation Strategies - MDPI, accessed on July 17, 2025, <https://www.mdpi.com/2413-4155/6/1/3>
10. Amazon's sexist hiring algorithm could still be better than a human - IMD Business School, accessed on July 17, 2025, <https://www.imd.org/research-knowledge/digital/articles/amazons-sexist-hiring-algorithm-could-still-be-better-than-a-human/>
11. Algorithmic bias, data ethics, and governance: Ensuring fairness, transparency and compliance in AI-powered business analytics applications, accessed on July 17, 2025, [https://journalwjarr.com/sites/default/files/fulltext\\_pdf/WJARR-2025-0571.pdf](https://journalwjarr.com/sites/default/files/fulltext_pdf/WJARR-2025-0571.pdf)
12. Legal Challenges in Regulating Algorithmic Bias and Discrimination - ResearchGate, accessed on July 17, 2025, [https://www.researchgate.net/publication/388122417\\_Legal\\_Challenges\\_in\\_Regulating\\_Algorithmic\\_Bias\\_and\\_Discrimination](https://www.researchgate.net/publication/388122417_Legal_Challenges_in_Regulating_Algorithmic_Bias_and_Discrimination)
13. How Artificial Intelligence Can Deepen Racial and Economic Inequities | ACLU, accessed on July 17, 2025, <https://www.aclu.org/news/privacy-technology/how-artificial-intelligence-can-deepen-racial-and-economic-inequities>
14. The Algorithmic Problem in Artificial Intelligence Governance | United Nations

- University, accessed on July 17, 2025,  
<https://unu.edu/article/algorithmic-problem-artificial-intelligence-governance>
15. Why Amazon's Automated Hiring Tool Discriminated Against Women | ACLU, accessed on July 17, 2025,  
<https://www.aclu.org/news/womens-rights/why-amazons-automated-hiring-tool-discriminated-against>
  16. The EEOC Puts Employers on Notice; The Use of AI in Hiring/Recruiting - Francis King Carey School of Law - The University of Maryland, Baltimore, accessed on July 17, 2025,  
<https://www.law.umaryland.edu/content/articles/name-660254-en.html>
  17. AI and Racial Bias in Legal Decision-Making: A Student Fellow Project, accessed on July 17, 2025,  
<https://clp.law.harvard.edu/knowledge-hub/insights/ai-and-racial-bias-in-legal-decision-making-a-student-fellow-project/>
  18. How We Analyzed the COMPAS Recidivism Algorithm - ProPublica, accessed on July 17, 2025,  
<https://www.propublica.org/article/how-we-analyzed-the-compas-recidivism-algorithm>
  19. Reprogramming Fairness: Affirmative Action in Algorithmic Criminal Sentencing, accessed on July 17, 2025,  
<https://hrlr.law.columbia.edu/hrlr-online/reprogramming-fairness-affirmative-action-in-algorithmic-criminal-sentencing/>
  20. The Age of Secrecy and Unfairness in Recidivism Prediction - Harvard Data Science Review, accessed on July 17, 2025,  
<https://hdsr.mitpress.mit.edu/pub/7z10o269>
  21. Injustice Ex Machina: Predictive Algorithms in Criminal Sentencing | UCLA Law Review, accessed on July 17, 2025,  
<https://www.uclalawreview.org/injustice-ex-machina-predictive-algorithms-in-criminal-sentencing/>
  22. The Ethics of AI in Surveillance and Privacy: Balancing Innovation with Human Rights, accessed on July 17, 2025,  
[https://www.researchgate.net/publication/390342528\\_The\\_Ethics\\_of\\_AI\\_in\\_Surveillance\\_and\\_Privacy\\_Balancing\\_Innovation\\_with\\_Human\\_Rights](https://www.researchgate.net/publication/390342528_The_Ethics_of_AI_in_Surveillance_and_Privacy_Balancing_Innovation_with_Human_Rights)
  23. (PDF) Regulating Autonomous AI: Legal Perspectives on Accountability and Liability, accessed on July 17, 2025,  
[https://www.researchgate.net/publication/392700339\\_Regulating\\_Autonomous\\_AI\\_Legal\\_Perspectives\\_on\\_Accountability\\_and\\_Liability](https://www.researchgate.net/publication/392700339_Regulating_Autonomous_AI_Legal_Perspectives_on_Accountability_and_Liability)
  24. Meaningful Human Control over Autonomous Systems: A Philosophical Account - PMC, accessed on July 17, 2025,  
<https://pmc.ncbi.nlm.nih.gov/articles/PMC7806098/>
  25. AI pitfalls and what not to do: mitigating bias in AI - PMC, accessed on July 17, 2025,  
<https://pmc.ncbi.nlm.nih.gov/articles/PMC10546443/>
  26. What are the most effective techniques for reducing bias in AI models trained on imbalanced datasets? | ResearchGate, accessed on July 17, 2025,  
[https://www.researchgate.net/post/What\\_are\\_the\\_most\\_effective\\_techniques\\_for](https://www.researchgate.net/post/What_are_the_most_effective_techniques_for)

- [\\_reducing\\_bias\\_in\\_AI\\_models\\_trained\\_on\\_imbalanced\\_datasets](#)
27. Bias recognition and mitigation strategies in artificial intelligence healthcare applications - PMC - PubMed Central, accessed on July 17, 2025, <https://pmc.ncbi.nlm.nih.gov/articles/PMC11897215/>
  28. EU AI Act: first regulation on artificial intelligence | Topics - European Parliament, accessed on July 17, 2025, <https://www.europarl.europa.eu/topics/en/article/20230601STO93804/eu-ai-act-first-regulation-on-artificial-intelligence>
  29. Ethical AI: Addressing bias, fairness, and accountability in autonomous decision-making systems - World Journal of Advanced Research and Reviews, accessed on July 17, 2025, <https://wjarr.com/sites/default/files/WJARR-2024-2510.pdf>
  30. The Blame Game: Legal Liability of AI in Autonomous Decision-Making - Medium, accessed on July 17, 2025, <https://medium.com/@kyrinstitute/the-blame-game-legal-liability-of-ai-in-autonomous-decision-making-0e491362e307>
  31. (PDF) Ai-Driven Autonomous Vehicles And Legal Liability: Redefining Accountability In Human-Ai Collaborative Systems - ResearchGate, accessed on July 17, 2025, [https://www.researchgate.net/publication/393360180\\_Ai-Driven\\_Autonomous\\_Vehicles\\_And\\_Legal\\_Liability\\_Redefining\\_Accountability\\_In\\_Human-Ai\\_Collaborative\\_Systems](https://www.researchgate.net/publication/393360180_Ai-Driven_Autonomous_Vehicles_And_Legal_Liability_Redefining_Accountability_In_Human-Ai_Collaborative_Systems)
  32. Accident Liability Determination of Autonomous Driving Systems Based on Artificial Intelligence Technology and Its Impact on Public Mental Health - PubMed Central, accessed on July 17, 2025, <https://pmc.ncbi.nlm.nih.gov/articles/PMC9451969/>
  33. A Vision on What Explanations of Autonomous Systems are of Interest to Lawyers - Research Explorer - The University of Manchester, accessed on July 17, 2025, [https://research.manchester.ac.uk/files/279269958/Explanations\\_for\\_Lawyers.pdf](https://research.manchester.ac.uk/files/279269958/Explanations_for_Lawyers.pdf)
  34. AI, humans and loops. Being in the loop is only part of the... | by Pawel Rzeszucinski, PhD | Medium, accessed on July 17, 2025, [https://medium.com/@pawel.rzeszucinski\\_55101/ai-humans-and-loops-04ee67ac820b](https://medium.com/@pawel.rzeszucinski_55101/ai-humans-and-loops-04ee67ac820b)
  35. Human-in-the-loop or AI-in-the-loop? Automate or Collaborate? - arXiv, accessed on July 17, 2025, <https://arxiv.org/html/2412.14232v1>
  36. Artificial intelligence liability directive - European Parliament, accessed on July 17, 2025, [https://www.europarl.europa.eu/RegData/etudes/BRIE/2023/739342/EPRS\\_BRI\(2023\)739342\\_EN.pdf](https://www.europarl.europa.eu/RegData/etudes/BRIE/2023/739342/EPRS_BRI(2023)739342_EN.pdf)
  37. Post #3 The AI Ethics Landscape: Government, Academia, and Business Approaches, accessed on July 17, 2025, <https://www.ethics.harvard.edu/blog/blog-post-3-ai-ethics-landscape-government-academia-and-business-approaches>
  38. When AI Cars Fail: Legal Guide - Charbonnet Law Firm, accessed on July 17, 2025, <https://www.charbonnetlawfirm.com/car-accidents/when-ai-cars-fail-legal-guide>

- [/](#)
39. How Will Artificial Intelligence Affect Jobs 2025-2030 | Nexford University, accessed on July 17, 2025, <https://www.nexford.edu/insights/how-will-ai-affect-jobs>
  40. AI-Driven Worker Displacement Is a Serious Threat - Jacobin, accessed on July 17, 2025, <https://jacobin.com/2025/07/artificial-intelligence-worker-displacement-jobs>
  41. AI and Economic Displacement - Unaligned Newsletter, accessed on July 17, 2025, <https://www.unaligned.io/p/ai-and-economic-displacement>
  42. The A.I. Workforce Displacement Crisis: Why Traditional Retraining Models Are Failing - Observer, accessed on July 17, 2025, <https://observer.com/2025/07/ai-job-loss-us-education-retraining/>
  43. AI, Automation, and the Urgent Case for Universal Basic Income - Scott Santens, accessed on July 17, 2025, <https://www.scottsantens.com/ai-automation-and-the-urgent-case-for-universal-basic-income-ubi-forward-future/>
  44. The effects of AI on firms and workers - Brookings Institution, accessed on July 17, 2025, <https://www.brookings.edu/articles/the-effects-of-ai-on-firms-and-workers/>
  45. Artificial Intelligence Impact on Labor Markets - International Economic Development Council (IEDC), accessed on July 17, 2025, [https://www.iedconline.org/clientuploads/EDRP%20Logos/AI\\_Impact\\_on\\_Labor\\_Markets.pdf](https://www.iedconline.org/clientuploads/EDRP%20Logos/AI_Impact_on_Labor_Markets.pdf)
  46. [www.innopharmaeducation.com](https://www.innopharmaeducation.com/blog/the-impact-of-ai-on-job-roles-workforce-and-employment-what-you-need-to-know#:~:text=These%20new%20jobs%20will%20be,training%20and%20teaching%20AI%20systems.), accessed on July 17, 2025, <https://www.innopharmaeducation.com/blog/the-impact-of-ai-on-job-roles-workforce-and-employment-what-you-need-to-know#:~:text=These%20new%20jobs%20will%20be,training%20and%20teaching%20AI%20systems.>
  47. Jobs AI will create? Here's the World Economic Forum view, accessed on July 17, 2025, <https://www.weforum.org/stories/2023/09/jobs-ai-will-create/>
  48. Could Basic Income Be The Answer to Jobs Lost to Automation and AI? A Critical Exploration and Conceptual Policy Framework - ResearchGate, accessed on July 17, 2025, [https://www.researchgate.net/publication/391808171\\_Could\\_Basic\\_Income\\_Be\\_The\\_Answer\\_to\\_Jobs\\_Lost\\_to\\_Automation\\_and\\_AI\\_A\\_Critical\\_Exploration\\_and\\_Conceptual\\_Policy\\_Framework](https://www.researchgate.net/publication/391808171_Could_Basic_Income_Be_The_Answer_to_Jobs_Lost_to_Automation_and_AI_A_Critical_Exploration_and_Conceptual_Policy_Framework)
  49. Universal basic income as a new social contract for the age of AI - LSE Business Review, accessed on July 17, 2025, <https://blogs.lse.ac.uk/businessreview/2025/04/29/universal-basic-income-as-a-new-social-contract-for-the-age-of-ai-1/>
  50. Universal Basic Income in the Age of Automation: A Critical Exploration and Policy Framework - ResearchGate, accessed on July 17, 2025, [https://www.researchgate.net/publication/391803859\\_Universal\\_Basic\\_Income\\_in\\_the\\_Age\\_of\\_Automation\\_A\\_Critical\\_Exploration\\_and\\_Policy\\_Framework](https://www.researchgate.net/publication/391803859_Universal_Basic_Income_in_the_Age_of_Automation_A_Critical_Exploration_and_Policy_Framework)
  51. AI, universal basic income, and power: symbolic violence in the tech elite's narrative - PMC, accessed on July 17, 2025,

- <https://pmc.ncbi.nlm.nih.gov/articles/PMC11891208/>
52. An AI Capability Threshold for Rent-Funded Universal Basic Income in an AI-Automated Economy - arXiv, accessed on July 17, 2025, <https://arxiv.org/html/2505.18687v1>
  53. Ways to help workers suffering from AI-related job losses - Brookings Institution, accessed on July 17, 2025, <https://www.brookings.edu/articles/ways-to-help-workers-suffering-from-ai-related-job-losses/>
  54. AI & the retraining challenge - AGI Social Contract, accessed on July 17, 2025, <https://www.agisocialcontract.org/anthology/ai-and-the-retraining-challenge>
  55. 'Politically unacceptable, morally repugnant': UN chief calls for global ban on 'killer robots', accessed on July 17, 2025, <https://news.un.org/en/story/2025/05/1163256>
  56. Using AI on Campuses: Security Surveillance or Privacy Invasion? | Pulitzer Center, accessed on July 17, 2025, <https://pulitzercenter.org/stories/using-ai-campuses-security-surveillance-or-privacy-invasion>
  57. Complexities of AI Trends: Threats to Data Privacy Legal Compliance, accessed on July 17, 2025, <https://digitalcommons.law.seattleu.edu/cgi/viewcontent.cgi?article=1096&context=sjteil>
  58. What Is AI Security? - IBM, accessed on July 17, 2025, <https://www.ibm.com/think/topics/ai-security>
  59. Artificial Intelligence in Education: Striking a Balance between Innovation & Privacy - Edly, accessed on July 17, 2025, <https://edly.io/blog/artificial-intelligence-in-education-and-privacy-concerns/>
  60. 7 AI Privacy Violations (+What Can Your Business Learn) - Enzuzo, accessed on July 17, 2025, <https://www.enzuzo.com/blog/ai-privacy-violations>
  61. the interdependence of artificial intelligence and global media ethics - ResearchGate, accessed on July 17, 2025, [https://www.researchgate.net/publication/393249821\\_THE\\_INTERDEPENDENCE\\_OF\\_ARTIFICIAL\\_INTELLIGENCE\\_AND\\_GLOBAL\\_MEDIA\\_ETHICS](https://www.researchgate.net/publication/393249821_THE_INTERDEPENDENCE_OF_ARTIFICIAL_INTELLIGENCE_AND_GLOBAL_MEDIA_ETHICS)
  62. AI and the Risk of Consumer Harm | Federal Trade Commission, accessed on July 17, 2025, <https://www.ftc.gov/policy/advocacy-research/tech-at-ftc/2025/01/ai-risk-consumer-harm>
  63. Understanding GDPR and CCPA in the Context of AI Systems - Signity Solutions, accessed on July 17, 2025, <https://www.signitysolutions.com/blog/understanding-gdpr-and-ccpa>
  64. The impact of the General Data Protection Regulation (GDPR) on artificial intelligence - European Parliament, accessed on July 17, 2025, [https://www.europarl.europa.eu/RegData/etudes/STUD/2020/641530/EPRS\\_STU\(2020\)641530\\_EN.pdf](https://www.europarl.europa.eu/RegData/etudes/STUD/2020/641530/EPRS_STU(2020)641530_EN.pdf)
  65. AI Gets Personal: CCPA vs. GDPR on Automated Decision-Making, accessed on July 17, 2025, <https://btlj.org/2025/04/ccpa-vs-gdpr-on-automated-decision-making/>

66. PPFL: Privacy-Preserving Techniques in Federated Learning, accessed on July 17, 2025, <https://journalaiai.com/index.php/aiai/article/view/35>
67. Federated Learning for Privacy-Preserving AI: An In-Depth Exploration - DSS Blog, accessed on July 17, 2025, <https://roundtable.datascience.salon/federated-learning-for-privacy-preserving-ai-an-in-depth-exploration>
68. federated learning with local differential privacy: trade-offs, accessed on July 17, 2025, <https://eprint.iacr.org/2021/142.pdf>
69. (PDF) Privacy-Preserving Federated Learning with Differential Privacy: Trade-offs and Implementation Challenges - ResearchGate, accessed on July 17, 2025, [https://www.researchgate.net/publication/392599662\\_Privacy-Preserving\\_Federated\\_Learning\\_with\\_Differential\\_Privacy\\_Trade-offs\\_and\\_Implementation\\_Challenges](https://www.researchgate.net/publication/392599662_Privacy-Preserving_Federated_Learning_with_Differential_Privacy_Trade-offs_and_Implementation_Challenges)
70. (PDF) AI in Data Privacy and Security. - ResearchGate, accessed on July 17, 2025, [https://www.researchgate.net/publication/378288596\\_AI\\_in\\_Data\\_Privacy\\_and\\_Security](https://www.researchgate.net/publication/378288596_AI_in_Data_Privacy_and_Security)
71. A Hazard to Human Rights: Autonomous Weapons Systems and Digital Decision-Making, accessed on July 17, 2025, <https://www.hrw.org/report/2025/04/28/hazard-human-rights/autonomous-weapons-systems-and-digital-decision-making>
72. Autonomous Weapons Systems and Transparency: Towards an International Dialogue - Scholarship Archive, accessed on July 17, 2025, [https://scholarship.law.columbia.edu/context/faculty\\_scholarship/article/5340/viewcontent/Knuckey\\_Autonomous\\_Weapons\\_Systems\\_and\\_Transparency.pdf](https://scholarship.law.columbia.edu/context/faculty_scholarship/article/5340/viewcontent/Knuckey_Autonomous_Weapons_Systems_and_Transparency.pdf)
73. International Discussions Concerning Lethal Autonomous Weapon Systems - Congress.gov, accessed on July 17, 2025, <https://www.congress.gov/crs-product/IF11294>
74. UN: Start Talks on Treaty to Ban 'Killer Robots' | Human Rights Watch, accessed on July 17, 2025, <https://www.hrw.org/news/2025/05/21/un-start-talks-treaty-ban-killer-robots>
75. Lethal Autonomous Weapons Systems & International Law: Growing Momentum Towards a New International Treaty | ASIL, accessed on July 17, 2025, <https://www.asil.org/insights/volume/29/issue/1>
76. The ethical implications of AI in warfare - Queen Mary University of London, accessed on July 17, 2025, <https://www.qmul.ac.uk/research/featured-research/the-ethical-implications-of-ai-in-warfare/>
77. What Role Do Ethical Frameworks Play in AI? - Lifestyle → Sustainability Directory, accessed on July 17, 2025, <https://lifestyle.sustainability-directory.com/question/what-role-do-ethical-frameworks-play-in-ai/>
78. Ethical and Responsible Use of AI for Students - CSU AI Commons, accessed on July 17, 2025, <https://genai.calstate.edu/communities/students/ethical-and-responsible-use-ai-students>



79. AI4People—An Ethical Framework for a Good AI Society: Opportunities, Risks, Principles, and Recommendations - PubMed Central, accessed on July 17, 2025, <https://pmc.ncbi.nlm.nih.gov/articles/PMC6404626/>
80. Responsible AI Systems: Who are the Stakeholders? - Open Research Online, accessed on July 17, 2025, <https://oro.open.ac.uk/84505/1/84505VOR.pdf>